



NEW YORK UNIVERSITY






Deep Learning

Kyunghyun Cho, Alfredo Canziani, Yann LeCun
NYU - Courant Institute & Center for Data Science

Deep Learning, NYU Fall 2020

Inspiration for Deep Learning: The Brain!

- ▶ McCulloch & Pitts (1943): networks of binary neurons can do logic
- ▶ Donald Hebb (1947): Hebbian synaptic plasticity
- ▶ Norbert Wiener (1948): cybernetics, optimal filter, feedback, autopoiesis, auto-organization.
- ▶ Frank Rosenblatt (1957): Perceptron
- ▶ Hubel & Wiesel (1960s): visual cortex architecture





Supervised Learning



- ▶ Training a machine by showing examples instead of programming it
- ▶ When the output is wrong, tweak the parameters of the machine

- ▶ Works well for:


- ▶ Speech → words
- ▶ Image → categories
- ▶ Portrait → name
- ▶ Photo → caption
- ▶ Text → topic
- ▶



CAR




PLANE




Supervised Learning goes back to the Perceptron & Adaline

► The McCulloch-Pitts Binary Neuron




- Perceptron: weights are motorized potentiometers
- Adaline: Weights are electrochemical “memistors”



$$y = \text{sign}\left(\sum_{i=1}^N W_i X_i + b\right)$$




<https://youtu.be/X1G2g3SiCwU>



The Standard Paradigm of Pattern Recognition

- ▶ ...since the 1960s
- ▶ ...and “traditional” Machine Learning
- ▶ until the “Deep Learning Revolution” (circa 2012)



Multilayer Neural Nets and Deep Learning

► Traditional Machine Learning



► Deep Learning



Parameterized Model

► Parameterized model

$$\bar{y} = G(x, w)$$

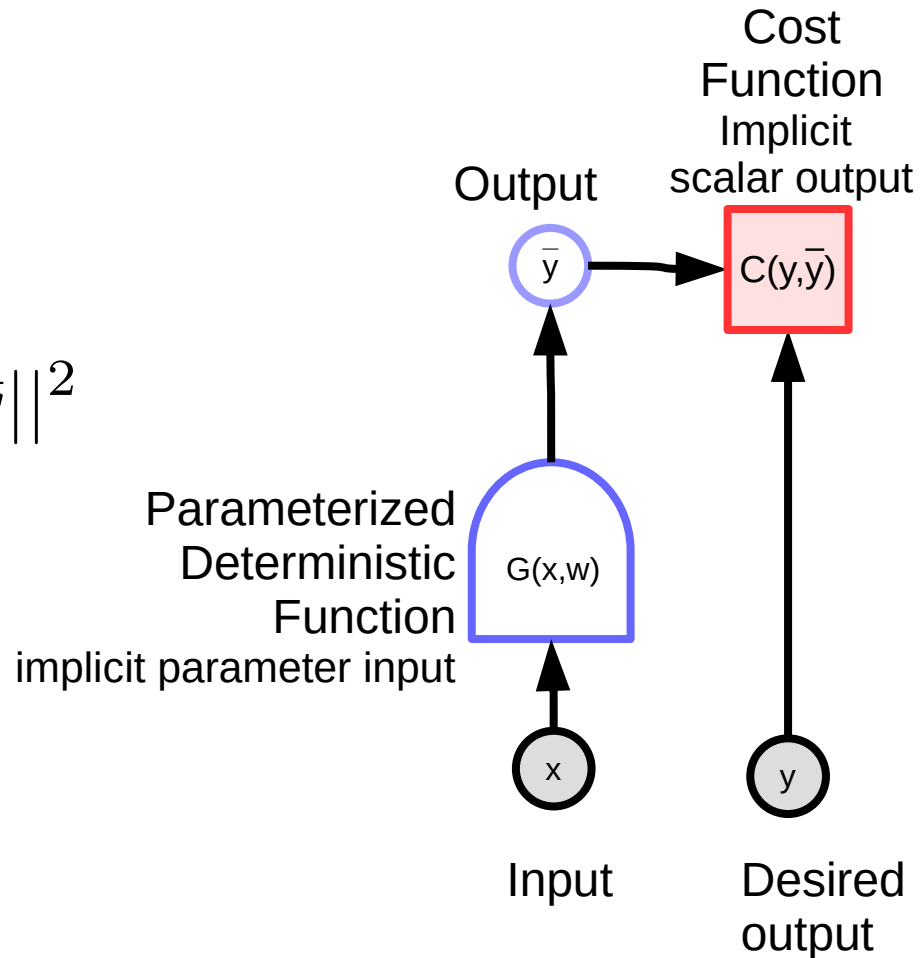
► Example: linear regression

$$\bar{y} = \sum_i w_i x_i \quad C(y, \bar{y}) = ||y - \bar{y}||^2$$

► Example: Nearest neighbor:

$$\bar{y} = \operatorname{argmin}_k ||x - w_{k,.}||^2$$

► Computing function G may involve complicated algorithms



Block diagram notations for computation graphs

▶ Variables (tensor, scalar, continuous, discrete...)




▶ Observed: input, desired output...



▶ Computed variable: outputs of deterministic functions


▶ Deterministic function



▶ Multiple inputs and outputs (tensors, scalars,...)

▶ Implicit parameter variable (here: w)

▶ Scalar-valued function (implicit output)



▶ Single scalar output (implicit)

▶ used mostly for cost functions

Loss function, average loss.

► Simple per-sample loss function


$$L(x, y, w) = C(y, G(x, w))$$

► A set of samples

$$S = \{(x[p], y[p]) \mid p = 0 \dots P - 1\}$$

► Average loss over the set

$$\mathcal{L}(S, w) = \frac{1}{P} \sum_{(x,y)} L(x, y, w) = \frac{1}{P} \sum_{p=0}^{P-1} L(x[p], y[p], w)$$



Supervised Machine Learning = Function Optimization



Function with
adjustable parameters




Objective
Function

traffic light: -1

■ It's like walking in the mountains in a fog and following the direction of steepest descent to reach the village in the valley

■ But each sample gives us a noisy estimate of the direction. So our path is a bit random.



$$W_i \leftarrow W_i - \eta \frac{\partial L(W, X)}{\partial W_i}$$

Gradient Descent

► Full (batch) gradient

$$w \leftarrow w - \eta \frac{\partial \mathcal{L}(S, w)}{\partial w}$$

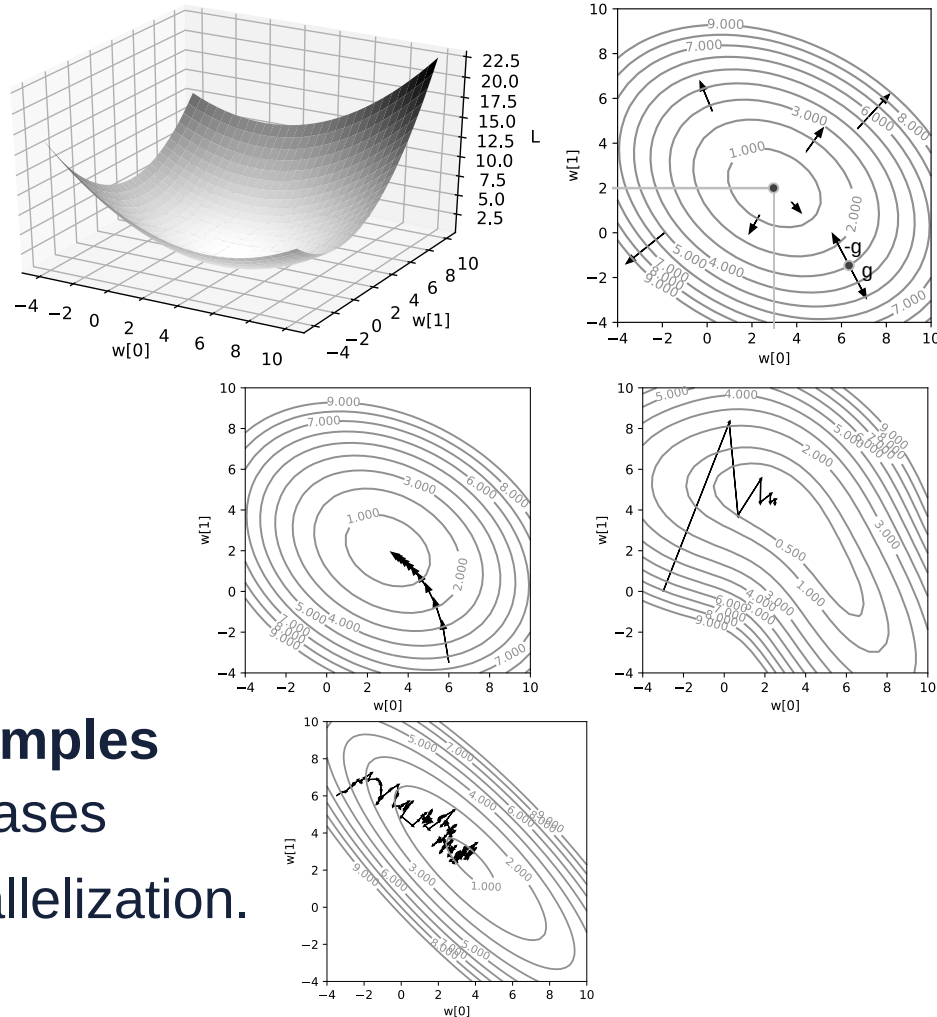
► Stochastic Gradient (SGD)

- Pick a p in $0 \dots P-1$, then update w :

$$w \leftarrow w - \eta \frac{\partial L(x[p], y[p], w)}{\partial w}$$


► SGD exploits the redundancy in the samples

- It goes faster than full gradient in most cases
- In practice, we use mini-batches for parallelization.





(Deep) Multi-Layer Neural Nets

- Multiple Layers of **simple units**
- Each units computes a **weighted sum** of its inputs
- Weighted sum is passed through a **non-linear** function
- The learning algorithm changes the **weights**




Ceci est une voiture



Traditional Neural Net


- ▶ **Stacked linear and non-linear functional blocks**
 - ▶ Weighted sums, matrix-vector product
 - ▶ Point-wise non-linearities (e.g. ReLu, tanh,)



Traditional Neural Net

► Stacked linear and non-linear functional blocks


$$s[i] = \sum_{j \in \text{UP}(i)} w[i, j] \cdot z[j] \quad z[i] = f(s[i])$$



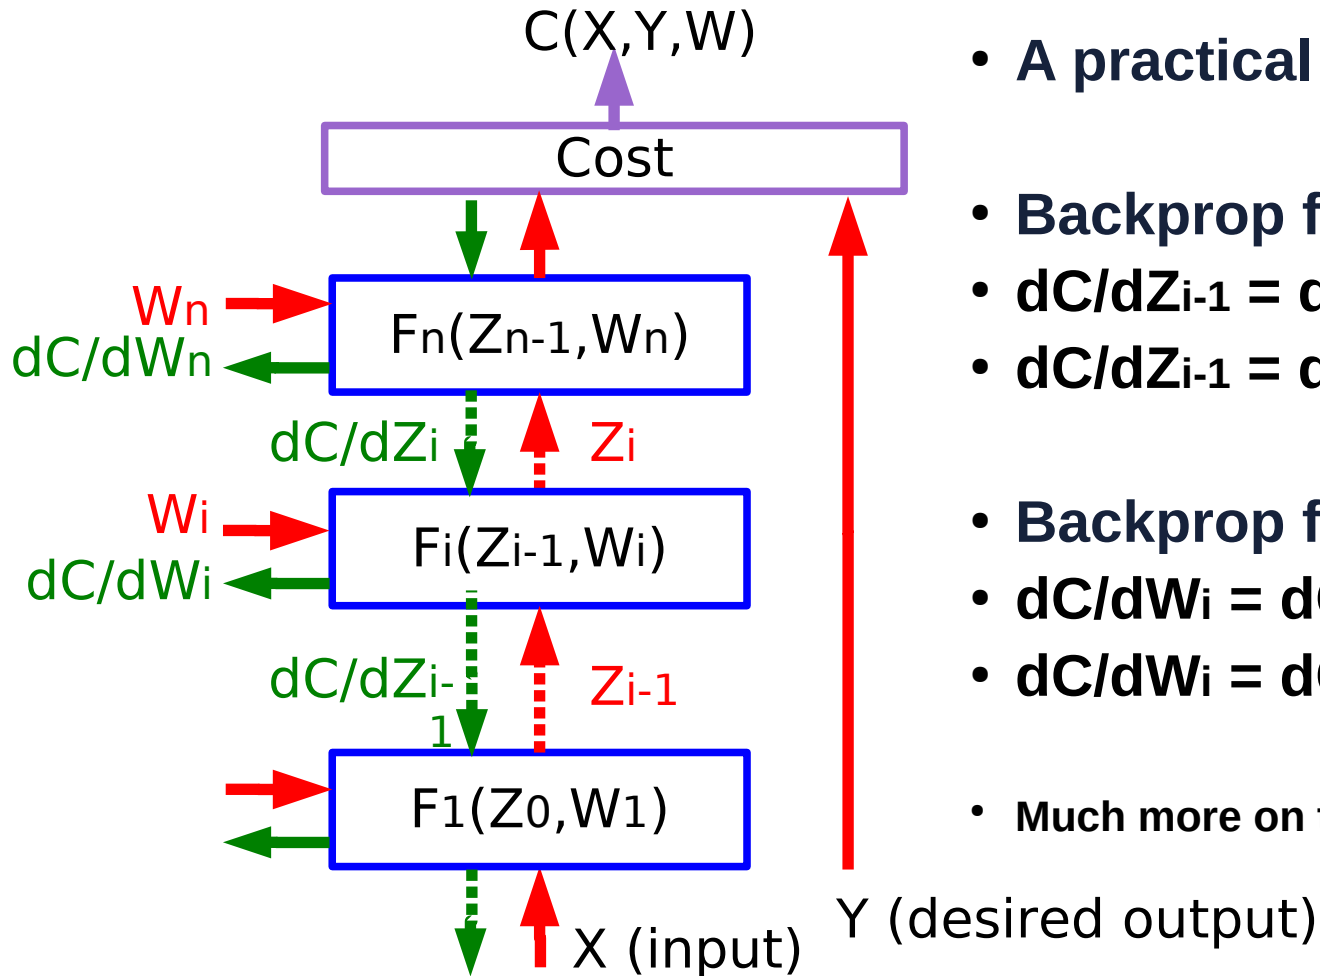
Block Diagram of a Traditional Neural Net

► linear blocks $s_{k+1} = w_k z_k$

► Non-linear blocks $z_k = h(s_k)$



Computing Gradients by Back-Propagation



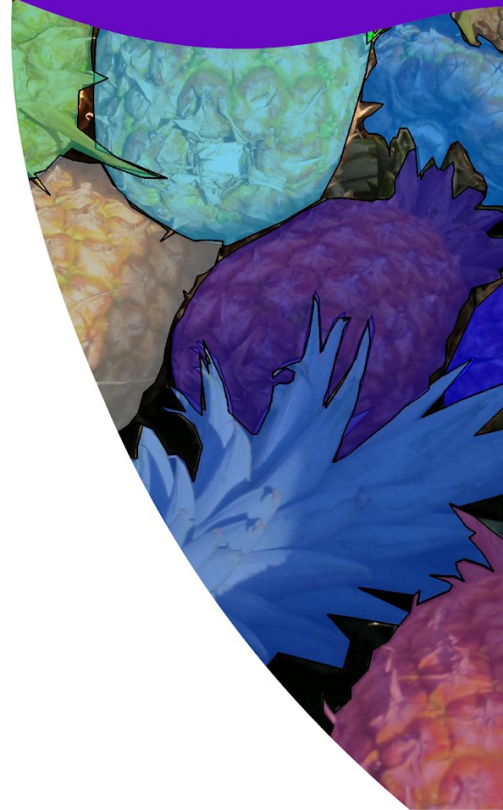
- A practical Application of Chain Rule
- Backprop for the state gradients:
- $dC/dZ_{i-1} = dC/dZ_i \cdot dZ_i/dZ_{i-1}$
- $dC/dZ_{i-1} = dC/dZ_i \cdot dF_i(Z_{i-1}, W_i)/dZ_{i-1}$
- Backprop for the weight gradients:
- $dC/dW_i = dC/dZ_i \cdot dZ_i/dW_i$
- $dC/dW_i = dC/dZ_i \cdot dF_i(Z_{i-1}, W_i)/dW_i$
- Much more on this later.....



NEW YORK UNIVERSITY

Learning Representations

What are good representations?
Why do networks need to be deep?



Deep Learning is about Learning Representations

► Traditional Machine Learning

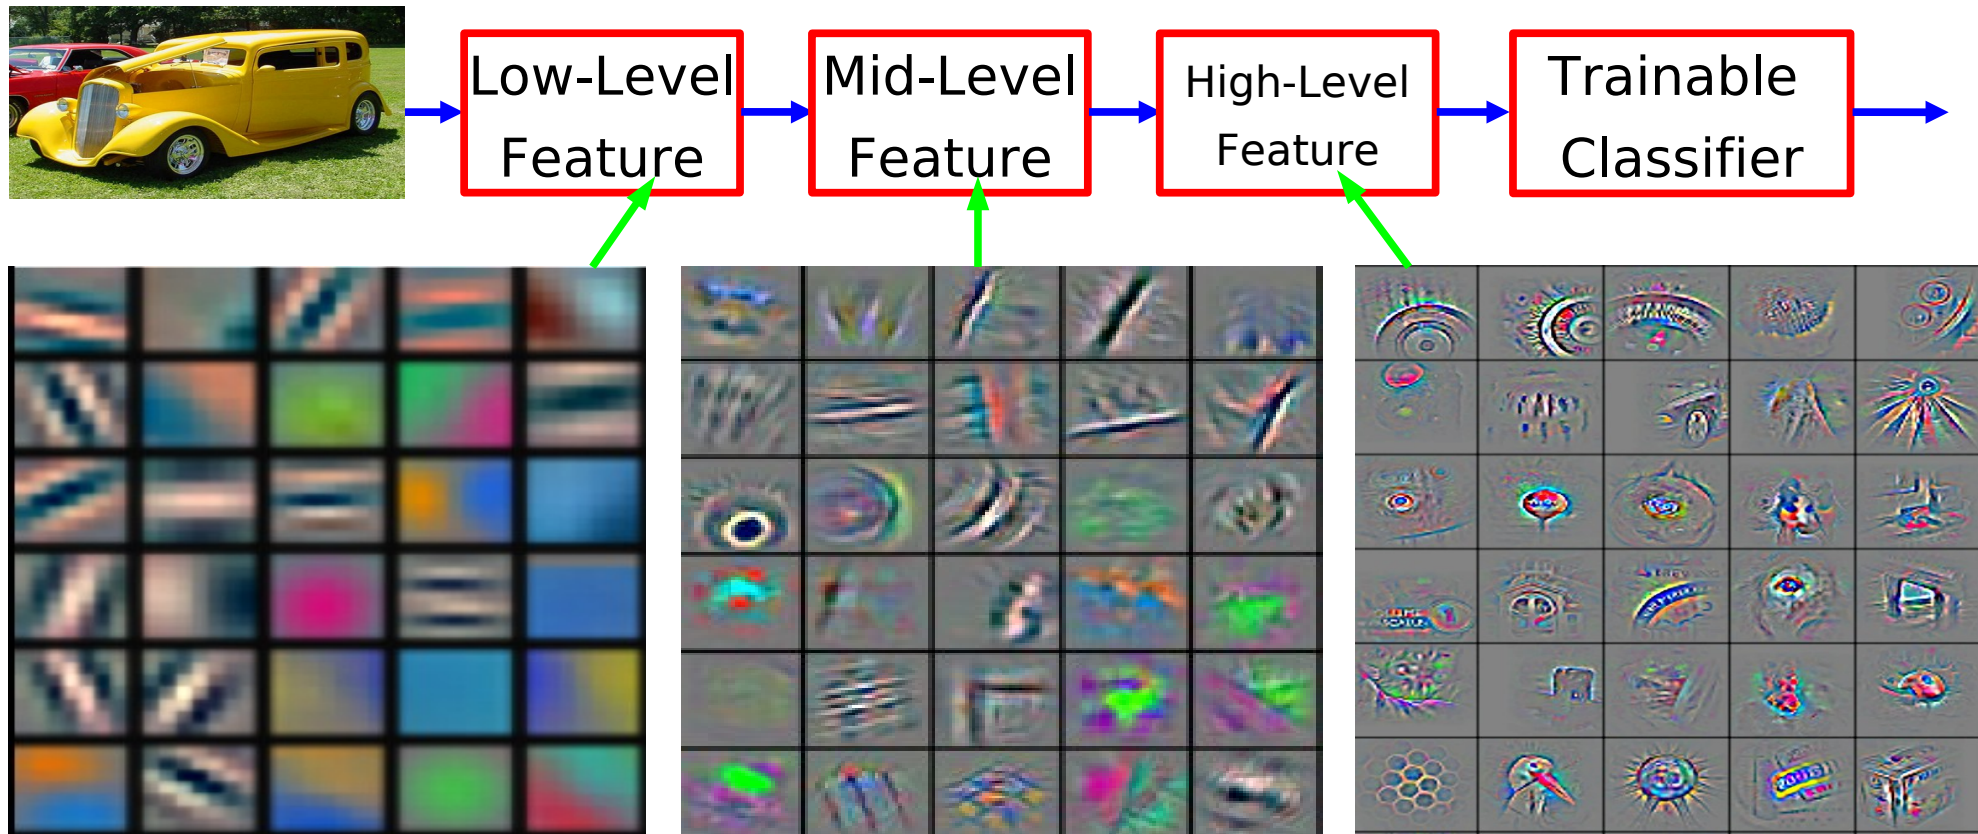


► Deep Learning



Multilayer Architectures == Compositional Structure of Data

■ Natural data is compositional => it is efficiently representable hierarchically



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

Multilayer Architecture == Hierarchical representation

- Hierarchy of representations with increasing level of abstraction

- Each stage is a kind of trainable feature transform

- Image recognition


 - ▶ Pixel → edge → texton → motif → part → object

- Text

 - ▶ Character → word → word group → clause → sentence → story

- Speech

 - ▶ Sample → spectral band → sound → ... → phone → phoneme → word



Shallow networks are universal approximators!


■ SVMs and Kernel methods

- ▶ Layer1: kernels; layer2: linear
- ▶ The first layer is “trained” with the simplest unsupervised method ever devised: using the samples as templates for the kernel functions.

■ 2-layer neural nets

- ▶ Layer1: dot products + non-linear function; Layer2: linear

■ But few useful functions can be efficiently represented with only two layers of reasonable size.



Ideas for “generic” feature extraction

- ▶ **Basic principle:**

- ▶ expanding the dimension of the representation so that things are more likely to become linearly separable.

- ▶ **- space tiling**

- ▶ **- random projections**

- ▶ **- polynomial classifier (feature cross-products)**

- ▶ **- radial basis functions**

- ▶ **- kernel machines**

Do we really need deep architectures?

 **Theoretician's dilemma:** “We can approximate any function as close as we want with shallow architecture. Why would we need deep ones?”

$$y = \sum_{i=1}^P \alpha_i K(X, X^i) \qquad y = F(W^1 . F(W^0 . X))$$


▶ kernel machines (and 2-layer neural nets) are “universal”.

 **Deep learning machines**

$$y = F(W^K . F(W^{K-1} . F(\dots F(W^0 . X) \dots)))$$

 **Deep machines are more efficient for representing certain classes of functions**, particularly those involved in visual recognition

▶ they can represent more complex functions with less “hardware”

 **We need an efficient parameterization of the class of functions that are useful for “AI” tasks (vision, audition, NLP...)**

Why would deep architectures be more efficient?

[Bengio & LeCun 2007 “Scaling Learning Algorithms Towards AI”]

A deep architecture trades space for time (or breadth for depth)

- ▶ more layers (more sequential computation),
- ▶ but less hardware (less parallel computation).

Example1: N-bit parity

- ▶ requires $N-1$ XOR gates in a tree of depth $\log(N)$.
- ▶ Even easier if we use threshold gates
- ▶ requires an exponential number of gates if we restrict ourselves to 2 layers (DNF formula with exponential number of minterms).

Example2: circuit for addition of 2 N-bit binary numbers

- ▶ Requires $O(N)$ gates, and $O(N)$ layers using N one-bit adders with ripple carry propagation.
- ▶ Requires lots of gates (some polynomial in N) if we restrict ourselves to two layers (e.g. Disjunctive Normal Form).
- ▶ Bad news: almost all boolean functions have a DNF formula with an exponential number of minterms $O(2^N)$